



Welcome to the Introduction to NoSQL course!



You are about to dive into to the exciting
world outside relational databases... 

About the course

- The course focuses on the diverse family of non-relational database management system, referred to commonly as **NoSQL database management systems**
- During this course, we will learn the principles of NoSQL DBMSs, how they differ from relational database management systems and what are their benefits and challenges
- We will be focusing mostly on a widely adopted NoSQL DBMS called **MongoDB**
- We will be covering MongoDB data management principles, data model, database operations and database programming in Python
- The prerequisite for the course is **Data Management and Databases and Python Programming** courses

Learning objectives

- The course has the following learning objectives:
 - Explain the main characteristics of NoSQL databases
 - Explain the main differences between NoSQL and relational databases
 - Provide examples of the main application areas of NoSQL databases
 - Explain the principles of the data model in the MongoDB database management system
 - Install the MongoDB database management system and the MongoDB Compass tool
 - Create a database in the MongoDB database management system
 - Perform database operations in the MongoDB database
 - Implement a simple Python application that operates on a MongoDB database

Materials and exercises

- The course material and exercises consists of **four parts**, which should be completed in the specified order:
 - The first part is **introduction to NoSQL DBMSs** with a multiple choice Moodle exam
 - The second part is **introduction to MongoDB DBMS** with a multiple choice Moodle exam and some practical exercises related to setting up the MongoDB DBMS
 - The third part covers the **MongoDB database operations** and data model with some practical exercises related to database queries
 - The fourth part covers **database programming in Python** with a small programming project exercise
- Course work submissions are in the Moodle area, but otherwise materials and exercises are in GitHub (link in the Moodle area)

Schedule and assessment

- The course schedule is very flexible, there are only **two deadlines**:
 - Passing the Moodle exam of the first part **before Thursday 30.10. at 23:59**
 - Finishing everything else **before Sunday 14.12. at 23:59.**
- The course assessment is either **pass or fail**. A passing grade requires passing the two Moodle exams and completing the other exercises in accordance with their requirements
- If your submission has issues, the teacher will request changes and you have a chance to resubmit your work

Brief history of NoSQL databases

- In the early 2000s, as the internet expanded, major companies such as Google and Amazon began to face unprecedented challenges related to data volume, velocity, and variety
- Traditional **relational database management systems** (RDBMSs) struggled to meet these demands, particularly in terms of scalability, flexibility, and performance in distributed environments
- To address these limitations, a new class of highly scalable systems, known as **NoSQL DBMSs**, began to emerge
- Since the 2010s, the NoSQL ecosystem has grown rapidly, with many NoSQL databases establishing themselves among the most widely adopted DBMSs in the world

What are NoSQL DBMSs?

"When people use the term "NoSQL database", they typically use it to refer to any non-relational database. Some say the term "NoSQL" stands for "non-SQL" while others say it stands for "not only SQL"."

- The term "NoSQL" is commonly referred to the large family of DBMSs that **are not relational**
- This means that the NoSQL DBMS's data model **is not based on relations and their relationships**
- Instead there are multiple different ways to structure data depending on the NoSQL DBMS, such as JSON documents, key-value pairs and graphs
- NoSQL DBMSs provide **solutions for a wide range of data requirements**. For instance, if we need extremely fast key-value operations without the overhead of managing complex relationships between data entities
- Next, let's have a look at three different NoSQL DBMSs and what's their data model approach: **MongoDB, Redis and Neo4j**

MongoDB – Document database, flexible, scalable

- MongoDB is an example of **document-based** NoSQL DBMS
- The data is stored in **collections** which contain JSON-like **documents**:

```
{  
  "_id": "68c3b961da0f719a26014775",  
  "title": "Dune",  
  "author": "Frank Herbert",  
  "published": 1965,  
  "genres": ["Science Fiction", "Adventure", "Politics"],  
  "copies": 7,  
  "location": {  
    "section": "Science Fiction",  
    "shelf": "S2"  
  }  
}
```

MongoDB – Document database, flexible, scalable

- Database operations are performed using the **MongoDB Query Language (MQL)**, which has a very JavaScript-like syntax:

```
db.books.find({ "author": "Frank Herbert" })
```

- The key benefits of MongoDB are its **flexibility** and **scalability**
- The JSON-like documents **don't have a fixed schema** (e.g. predefined field names and data-types), which makes it easy to evolve the structure
- MongoDB can handle large volumes of data and high traffic efficiently due **its ability to scale effortlessly** using techniques such as sharding
- These benefits make MongoDB a good choice for e.g. content management systems and real-time applications

Redis – Key-value database, fast, scalable


- Redis is an example of **key-value** NoSQL DBMS
- The data is stored as **key-value pairs**, meaning that we can store and access any kind of string data using string keys, e.g. unique identifiers:

```
SET teacher:h02680 "h02680;Kalle;Ilves;kalle.ilves@haaga-helia.fi"  
GET teacher:h02680
```

- The Redis data model and operations are quite limited, but the basic operations (e.g. `SET` and `GET`) are **extremely fast**
- Instead of accessing the data on disk, Redis **stores all data in main memory**
- Redis also saves snapshots of the dataset on disk, which means that the data is **persisted**
- Due to its fast key access operations, Redis is commonly used as a **cache layer** in many software systems

Neo4j – Graph database, relationships, efficient queries

- Neo4j is an example of a **graph-based** NoSQL DBMS
- The data is stored as **nodes** which have **relationships** with other nodes:



- Database operations are performed using the **Cypher** query language:

```
MATCH (a:Person)-[:ACTED_IN]->(m:Movie {title: "Inception"})  
RETURN a.name AS Actor;
```

- Neo4j is a good choice for all use-cases where the **data can be modelled as a graph**, e.g. social networks.

Why should we study NoSQL DBMSs?


"It is tempting, if the only tool you have is a hammer, to treat everything as if it were a nail."

- As we saw from these examples alone, the **NoSQL DBMS family offers a wide range of different kind of tools** suitable for different kind of use-cases
- For example, Redis offers a way better performance for key-value based operations than PostgreSQL or other RDBMSs
- MongoDB on the other hand offers database schema flexibility and built-in scalability
- So, studying NoSQL DBMSs offers us more tools to solve different kind of problems

The popularity of NoSQL DBMSs

- The "State of Database Survey" annually studies the usage of different database technologies. Here's a few highlights from the year 2023:
 - **MongoDB** and **Redis** NoSQL DBMSs are in the **top four most widely used DBMSs** right after PostgreSQL and MySQL
 - 74.4% percentage of respondents have used MongoDB and 73.3% have used Redis
- This indicates that **NoSQL DBMSs are widely acknowledged and used in the industry** despite not being as popular as certain relational DBMSs
- In web-development, different kind of **tech stacks containing MongoDB are widely adopted**
- MERN (MongoDB, Express, React, Node) and MEAN (MongoDB, Express, Angular, Node) are acronyms for such common tech stacks

HOW TO WRITE A CV



Leverage the NoSQL boom

Closing words & QA

- Thanks for joining the session and again, welcome to the course!
- Next steps:
 1. **Let's hear any questions you have**
 2. Once everything is clear, let's start the course! 
- Don't be afraid to contact me during the course regarding any questions you have