



Welcome to the Introduction to NoSQL course! 🙌

You are about to dive into the exciting world outside relational databases... 🙌

About the course

- This course focuses on the diverse and widely used family of non-relational database management systems, commonly called **NoSQL database management systems**
- We will learn the core principles of NoSQL DBMSs, how they differ from relational DBMSs, their application areas, and their benefits and challenges
- The main focus in the course is the **MongoDB** NoSQL DBMS, which is among the most popular DBMSs
- We will cover MongoDB data management principles, data modeling, database operations, and database programming in Python
- The prerequisites are the **Data Management and Databases** and **Python Programming** courses

Learning objectives

- After completing this course, you know how to...
 - Explain the main characteristics of NoSQL database management systems
 - Explain the main differences between NoSQL and relational DBMSs
 - Provide examples of the main application areas of NoSQL DBMSs
 - Explain the principles of the data model in the MongoDB DBMS
 - Install the MongoDB DBMS and the MongoDB Compass tool
 - Create a database and perform database operations with the MongoDB DBMS
 - Implement a simple Python application that operates on a MongoDB database

Materials and exercises

- The course material and exercises consist of **four parts**, which should be completed in order:
 - The first part is an **introduction to NoSQL DBMSs** with a multiple-choice Moodle exam
 - The second part is an **introduction to MongoDB DBMS** with a multiple-choice Moodle exam and practical setup exercises
 - The third part covers **MongoDB data modeling and database operations** with practical query exercises
 - The fourth part covers **database programming in Python** with a small programming project
- Coursework is submitted in Moodle, while the materials and exercises are in GitHub (link in Moodle)

Schedule and assessment

- The course schedule is flexible, and there are only **two deadlines**:
 - Passing the Moodle exam of the first part **before Thursday 29.10. at 23:59**
 - Finishing everything else **before Sunday 13.12. at 23:59**
- The course is assessed as **pass or fail**. A passing grade requires passing both Moodle exams and completing the other exercises according to their requirements
- If your submission has issues, the teacher will request changes and you have a chance to resubmit your work

Brief history of NoSQL databases

- As the internet expanded in the early 2000s, major companies such as Google and Amazon faced challenges in maintaining query performance with an ever-growing number of simultaneous users and highly varied data
- Traditional **relational DBMSs** struggled to meet these demands, particularly in terms of scalability, flexibility, and performance
- To address these limitations, a new class of **highly scalable** systems, known as **NoSQL DBMSs**, began to emerge from within these major companies
- One of the core design principles of many NoSQL DBMSs has been **horizontal scaling** (or scale-out), which distributes data storage and computation to multiple machines in the system
- Since the 2010s, **the NoSQL ecosystem has grown rapidly**, with many NoSQL DBMSs establishing themselves among the most widely adopted DBMSs in the world

What are NoSQL DBMSs?

"When people use the term "NoSQL database", they typically use it to refer to any non-relational database. Some say the term "NoSQL" stands for "non-SQL" while others say it stands for "not only SQL"."

- The term "NoSQL" commonly refers to the large family of DBMSs that **are not relational**, meaning the data model **is not based on tables, relationships, and join operations**
- Instead, data can be structured in multiple ways depending on the NoSQL DBMS, such as **JSON documents, key-value pairs, and graphs**
- Different NoSQL DBMSs provide **solutions for many different data requirements**, for example extremely fast key-value access without the overhead of managing complex table relationships
- Next, let's look at three widely used NoSQL DBMSs and their data model approaches:
MongoDB, Redis, and DynamoDB

MongoDB — Document database



- **MongoDB** is a **document database** that stores data in **collections** containing flexible, JSON-like **documents**:

```
{
  "_id": "68c3b961da0f719a26014775",
  "title": "Dune",
  "author": "Frank Herbert",
  "published": 1965,
  "genres": ["Science Fiction", "Adventure", "Politics"],
  "copies": 7,
  "ebook": false,
  "location": {
    "section": "Science Fiction",
    "shelf": "S2"
  }
}
```


MongoDB — Document database



- Database operations are performed using the **MongoDB Query Language** (MQL), which has a very JavaScript-like syntax:

```
db.books.find({ "author": "Frank Herbert" })
```

- The key benefits of MongoDB are its **flexibility** and **horizontal scalability**
- The JSON-like documents **don't have a fixed schema** (e.g. predefined field names and data types), which makes it easy to evolve the structure
- MongoDB can handle high data volumes and traffic efficiently due to **its ability to scale using sharding and replication**
- These benefits make MongoDB a good choice for content management systems (flexibility requirements due to varying data) and real-time applications (high scalability requirements due to high data volumes)

Redis — Key-value database




- **Redis** is a **key-value database** that stores data as **key-value pairs**, where a unique string key maps to a value in a flexible format
- The queries focus on key-based operations, e.g. accessing the value of a specific key:

```
SET teacher:h02680 "h02680;Kalle;Ilves;kalle.ilves@haaga-helia.fi"  
GET teacher:h02680
```

- Redis has a simpler data model than MongoDB, but its core operations (e.g., `SET` and `GET`) are **extremely fast due to its in-memory design**
- Data is stored in main memory for fast access, while optionally persisting it to disk for permanent storage
- Due to its fast key-based access, Redis is commonly used as a **secondary caching layer** to store frequently accessed, slow-to-query data, reducing latency and offloading the primary database

DynamoDB — Key-value and document database



- **Amazon DynamoDB** is a **key-value and document database** managed in AWS
- The data is stored in **tables** (different from relational DBMSs) that contain **items**, which are flexible JSON documents identified by a primary key
- Database operations are performed using the **DynamoDB API**:

```
{  
  "TableName": "Orders",  
  "KeyConditionExpression": "CustomerId = :id",  
  "ExpressionAttributeValues": {  
    ":id": { "S": "cus_7mbrsagJ7" }  
  },  
  "ProjectionExpression": "OrderId, Items, ShippingAddress"  
}
```

- DynamoDB's queries are primarily designed around **fast key-based access**, in contrast to MongoDB, which supports flexible and complex queries across many fields

Trade-offs in NoSQL DBMSs

- The benefits of NoSQL DBMSs come with drawbacks, and **trade-offs need to be made to gain certain benefits**
- **Horizontal scalability** offers performance and fault tolerance, but often trades away **strong consistency**. Written data is only "*eventually*" propagated across all machines in the system (eventual consistency) and until then, users may read **stale, outdated data**
- **Schema flexibility** makes it easy to evolve the data model, but can lead to **low data quality**, since no strict rules are enforced on the structure of the data (though enforcement is possible in some NoSQL DBMSs, such as MongoDB)
- In many scenarios, however, these trade-offs are reasonable and can still meet the data requirements

Why should we study NoSQL DBMSs?

"It is tempting, if the only tool you have is a hammer, to treat everything as if it were a nail."

- As we saw from these examples alone, the **NoSQL DBMS family offers a wide range of different kinds of tools** suitable for different kinds of use-cases
- For example, Redis often offers much better performance for key-value operations than relational RDBMSs
- MongoDB, on the other hand, offers database schema flexibility and built-in scalability
- Often, a system's architecture **consists of more than one DBMS**, for example MongoDB as the primary database and Redis as the cache layer
- Studying NoSQL DBMSs offers us more tools to solve different kinds of problems

The popularity of NoSQL DBMSs

- The "State of Database Survey" annually studies the usage of different database technologies. Here are a few highlights from the year 2023:
 - **MongoDB** and **Redis** NoSQL DBMSs are in the **top four most widely used DBMSs** right after PostgreSQL and MySQL
 - 74.4% of respondents have used MongoDB and 73.3% have used Redis
- This indicates that **NoSQL DBMSs are widely acknowledged and used in the industry** despite not being as popular as certain relational DBMSs
- In web development, different kinds of **tech stacks containing MongoDB are widely adopted**
- MERN (MongoDB, Express, React, Node) and MEAN (MongoDB, Express, Angular, Node) are acronyms for such common tech stacks

HOW TO WRITE A CV



Leverage the NoSQL boom

Closing words & QA

- Thanks for joining the session and again, welcome to the course!
- Next steps:
 1. **Let's hear any questions you have**
 2. Once everything is clear, let's start the course! 🧙
- Don't be afraid to contact me during the course regarding any questions you have